



NWEN 241

Getting closer to the system

Process Management

Winston Seah

School of Engineering and Computer Science
Victoria University of Wellington

Victoria
UNIVERSITY OF WELLINGTON

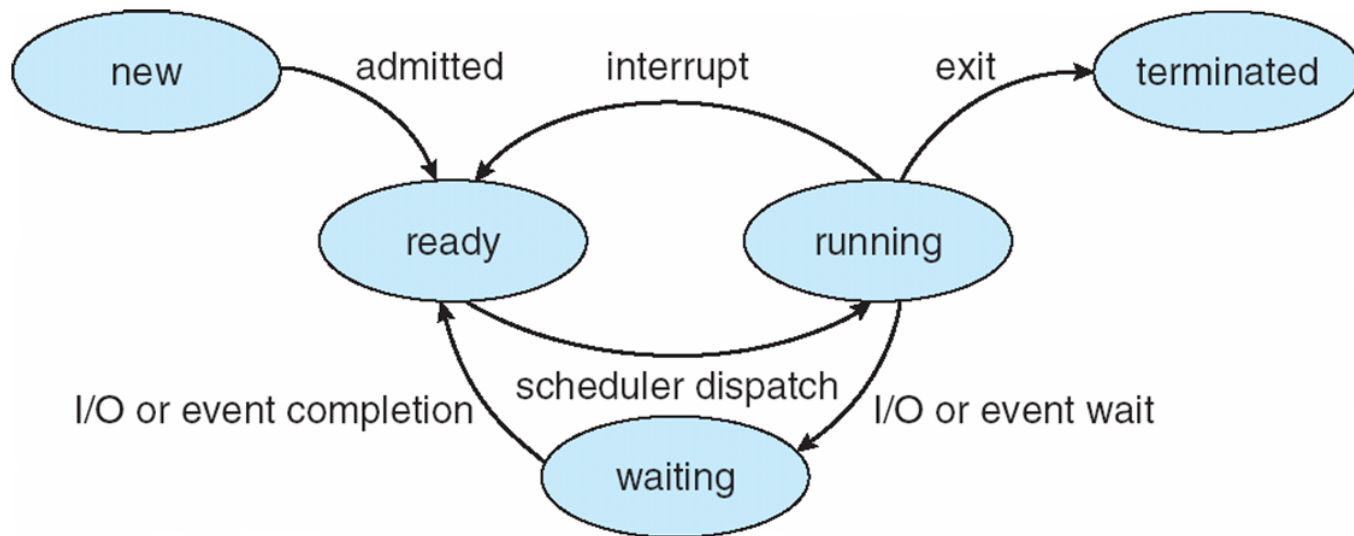
*Te Whare Wānanga
o te Ūpoko o te Ika a Māui*





CAPITAL CITY UNIVERSITY

Process vs Program

- Process – a program in execution
- A process includes (among other things):
 - program counter
 - data section
 - stack
- Process life-cycle



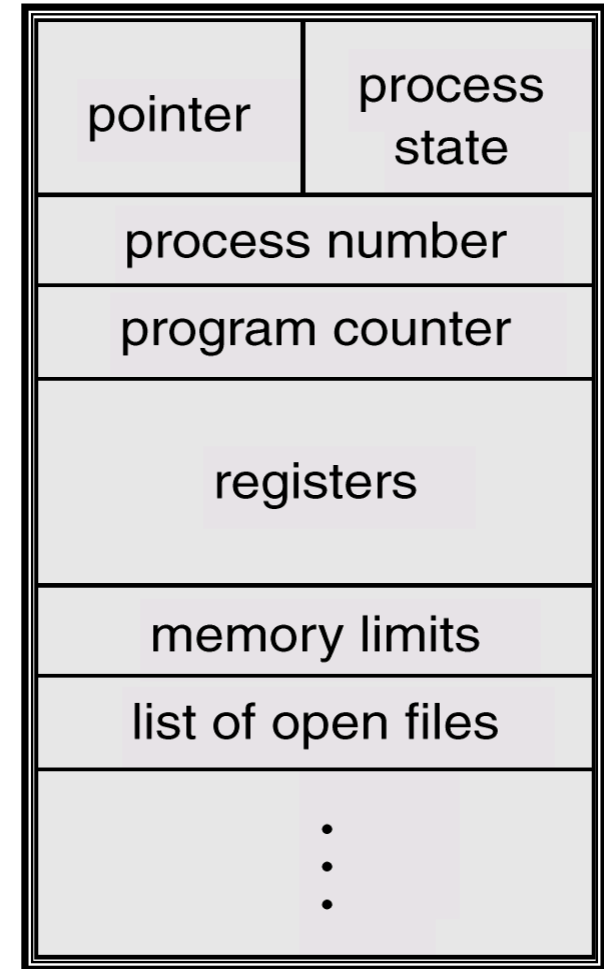
Process vs Program



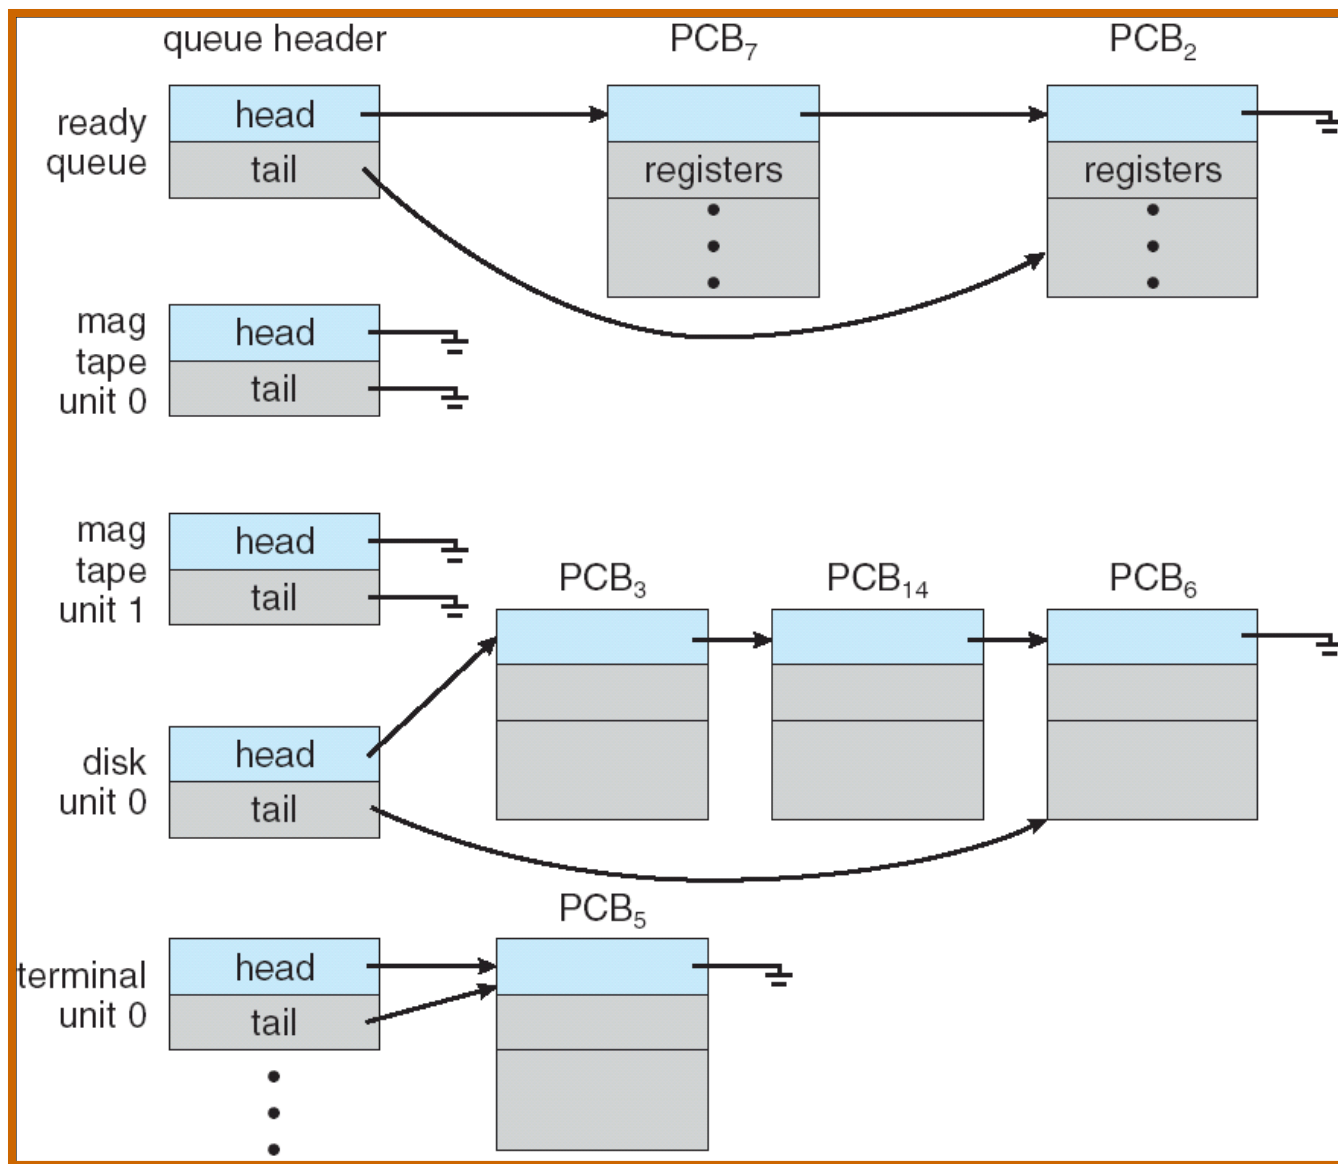
- Program is static, with the potential for execution
- Process is a program in execution and have a state
- One program can be executed several times and thus has several processes

Process Control Block (PCB)

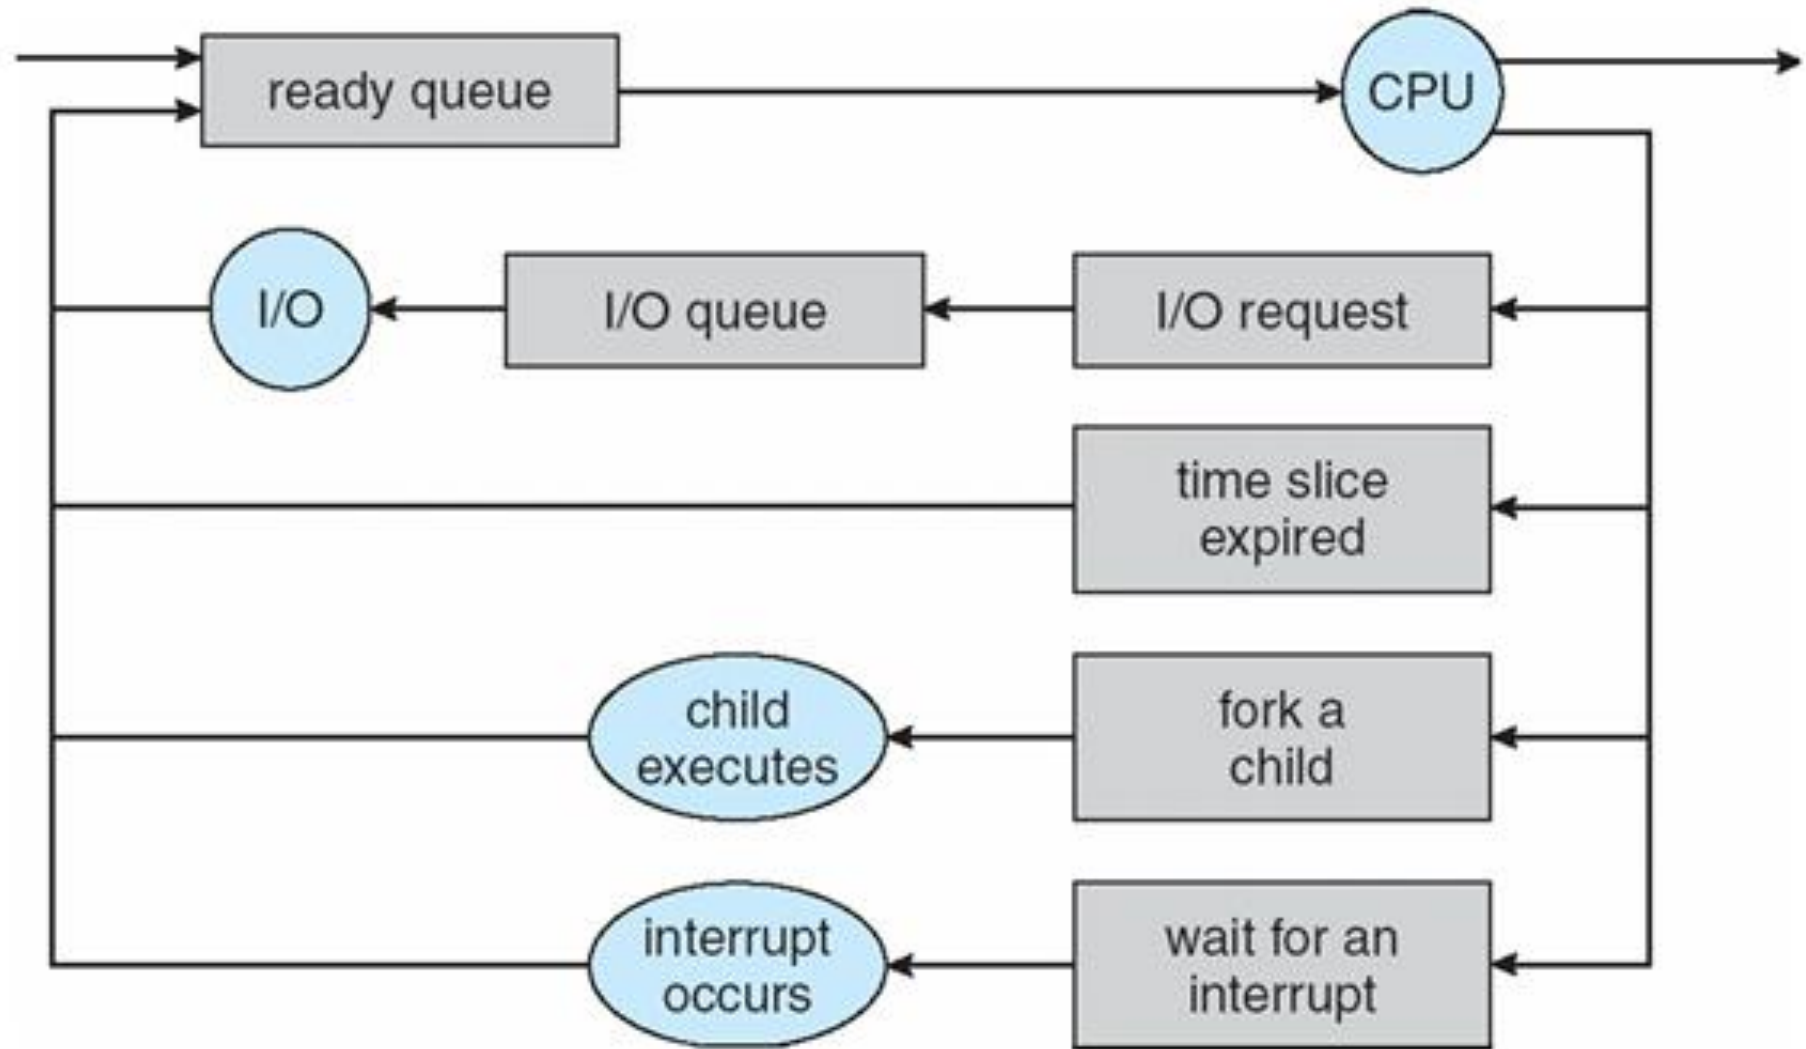
- Information associated with each process
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information
- A process is named using its process ID (PID) or process #
- Data is stored in a process control block (PCB)



Ready Queue and I/O Device Queues

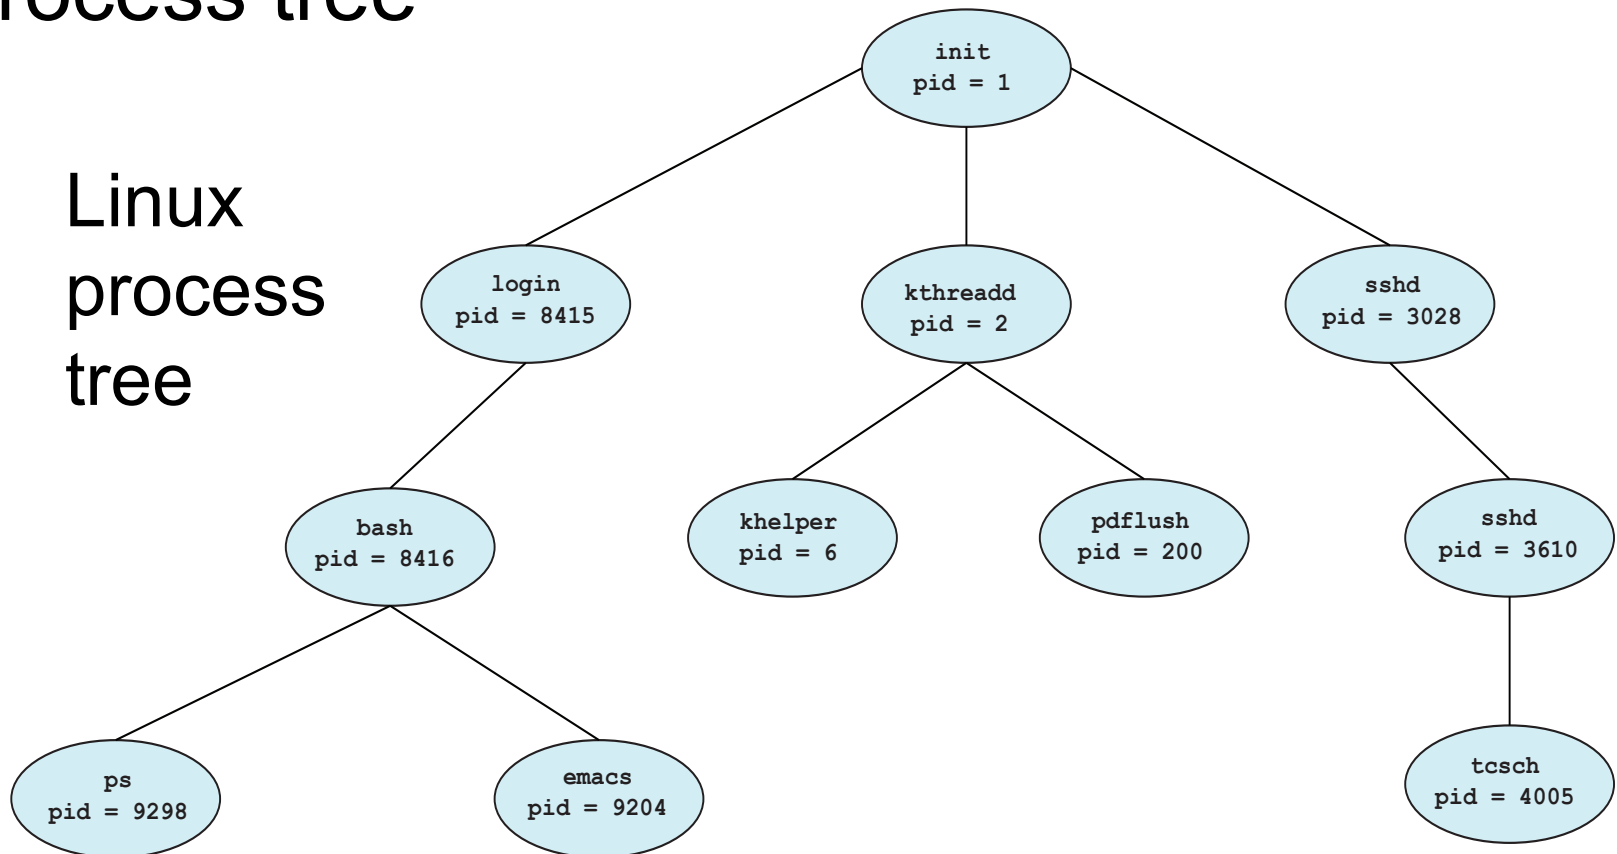


Process Scheduling / Switching



Process Management

- A process is created by another process, which, in turn create other processes → process tree




Process Management

- Parent and child process
 - In Linux, using “ps -f”, the PPID field is the parent
 - The first process is *init* having process ID 1
- After creating a child, the parent may either wait for it to finish or continue concurrently
- Process Management in C using System Calls
 - `fork()`
 - `exec()`
 - `wait()`
 - `exit()`

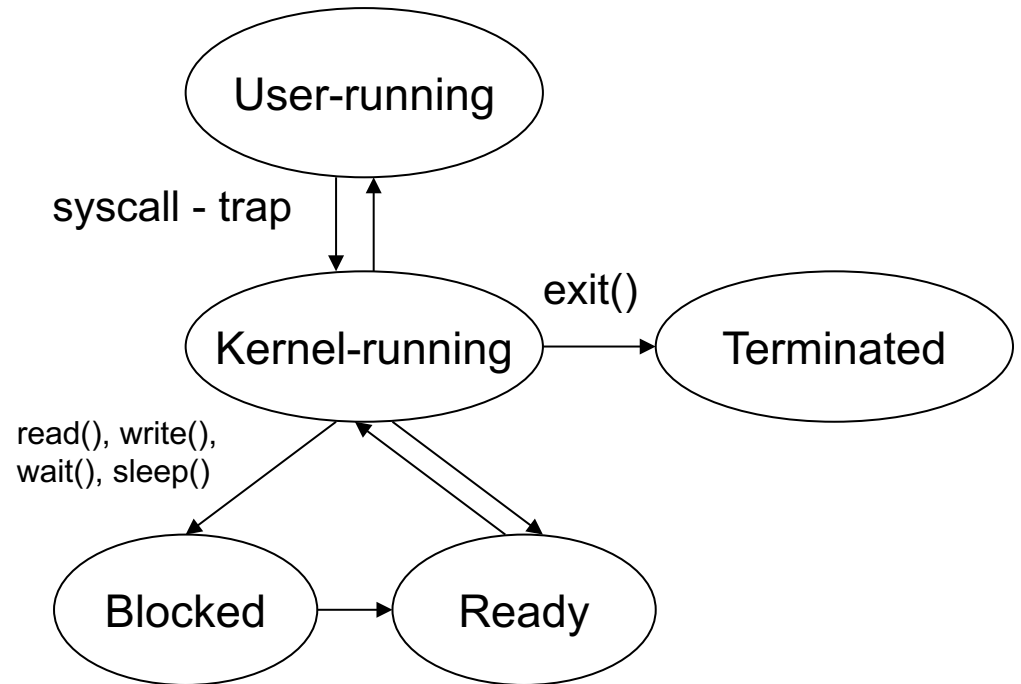

System Call

- A direct request to the operating system to do something on behalf of the program
- Typically programs are executed in **user** mode
- System call allows a switch from **user** mode to **kernel** mode




- The **kernel** is the core of the operating system for managing **processes**, **files**, **networking**, etc..

System Call Interface



Process Creation with `fork()` System Call

- A process calling `fork()` spawns a child process.
- After a successful `fork()` call, two copies of the original code will be running.
 - Parent process – *return value of `fork()` → child PID.*
 - New child process – *return value of `fork()` → 0.*
- `fork()` is called once, but returns twice!
- After `fork()` both the parent and the child are executing the same program.
- On error, `fork()` returns -1.



Consider a piece of program
(see examples latex):

```
...  
pid_t pid = fork();  
printf("PID: %d\n", pid);  
...
```

The parent will print:

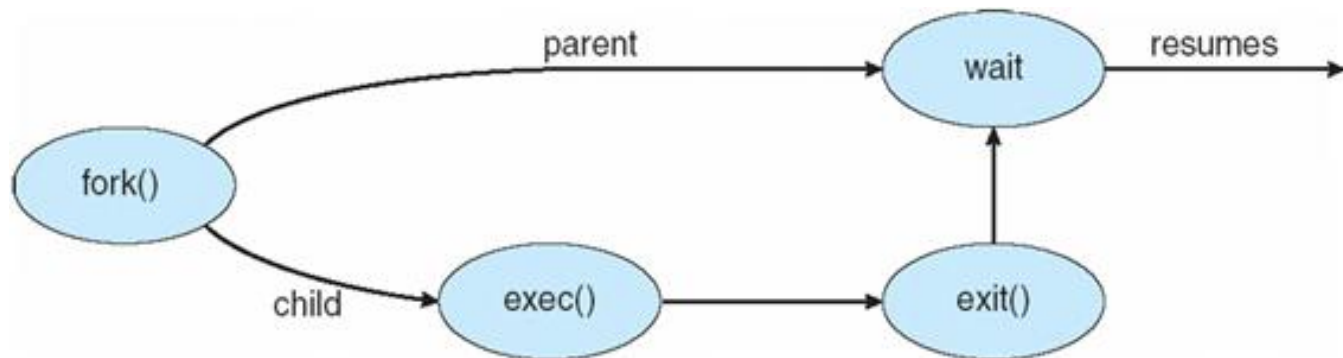
PID: 34

And the child will **always** print:

PID: 0

exec () System call (1)

- The **exec ()** call replaces a current process' image with a new one (i.e. loads a new program within current process).
- Upon success, **exec ()** **never** returns to the caller. If it does return, it means the call failed. Typical reasons are: non-existent file (bad path) or bad permissions.
- Arguments passed via **exec ()** appear in the **argv []** of the **main ()** function.

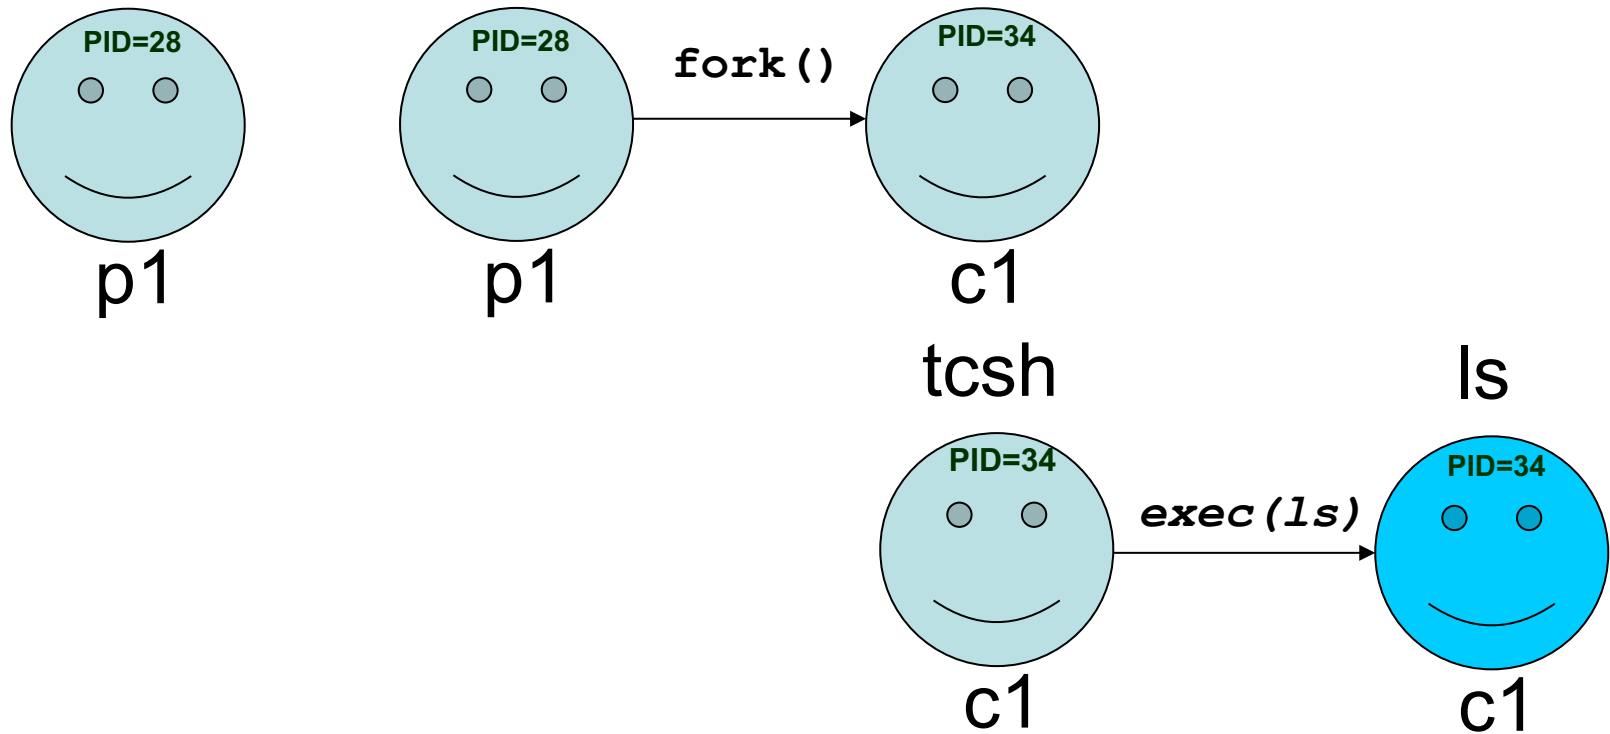


exec () System call (2)

- There is **no** system call specifically by the name **exec ()**. By **exec ()** we usually refer to a family of calls:
 - `int execl(char *path, char *arg, ...);`
 - `int execv(char *path, char *argv[]);`
 - `int execl(char *path, char *arg, ..., char *envp[]);`
 - `int execve(char *path, char *argv[], char *envp[]);`
 - `int execlp(char *file, char *arg, ...);`
 - `int execvp(char *file, char *argv[]);`
- The various options *l*, *v*, *e*, and *p* mean:
 - *l* : an argument list,
 - *v* : an argument vector,
 - *e* : an environment vector, and
 - *p* : a search path.

fork () and exec () together

- Often after doing `fork ()` we want to load a new program into the child.
- Most common e.g. a shell.



Example of forking separate processes

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

wait() Call System

- Forces the parent to suspend execution, i.e. wait for its children or a specific child to die (*terminate* is more appropriate terminology, but a bit less common).

```
pid_t wait(int *status);
```

- `waitpid()` waits for the child with specific PID.

```
pid_t waitpid(pid_t pid, int *status,  
              int options);
```

- The status, if not NULL, stores exit information of the child, which can be analyzed by the parent.
- The return value is:
 - PID of the exited process, if no error
 - (-1) if an error has happened

`exit()` System Call

- Gracefully terminates process execution, meaning it does clean up and release of resources, and puts the process into the **zombie** state → *terminated but still waiting for parent process to read its exit status*.
- By calling `wait()`, the parent cleans up all its zombie children.
- `exit()` specifies a return value from the program, which a parent process might want to examine as well as status of the dead process.
- `_exit()` call is another possibility of quick death without cleanup.

Example of wait() and exit()

```
#include <stdio.h>
#include <stdlib.h>

main()
{
    int pid; int rv;

    pid=fork();
    switch(pid){
        case -1:
            printf("Error -- Something went wrong with fork()\n");
            exit(1); // parent exits
        case 0:
            printf("CHILD: This is the child process!\n");
            printf("CHILD: My PID is %d\n", getpid());
            printf("CHILD: My parent's PID is %d\n", getppid());
            printf("CHILD: Enter my exit status: ");
            scanf(" %d", &rv);
            printf("CHILD: I'm outta here!\n");
            exit(rv);
        default:
            printf("PARENT: This is the parent process!\n");
            printf("PARENT: My PID is %d\n", getpid());
            printf("PARENT: My child's PID is %d\n", pid);
            printf("PARENT: I'm now waiting for my child to exit()...\n");
            wait(&rv);
            printf("PARENT: I'm outta here!\n");
    }
}
```

More about `wait()` and `exit()`

- Should not interpret the status value of system call `wait(&status)` literally. If `&status` is not NULL, `wait()` stores status information in the `int` to which it points.
- Value returned by `exit(&status)` is moved to 2nd byte and 1st (lowest) byte is used to store the status information.
- In previous example:

```
scanf(" %d", &rv); // if value of x is entered
```


```
...
```

```
exit(rv);
```



```
...
```

```
wait(&rv); // the rv contents will be x left shift by  
           // 8 bits and additional status written into  
           // lowest 8 bit
```



More about `wait()` and `exit()`

- This `status` integer can be inspected with macros (which take the integer itself as an argument, **not** a pointer to it):
 - `WIFEXITED(status)`
 - `WEXITSTATUS(status)` This macro should be employed only if `WIFEXITED` returned true.
 - `WIFSIGNALED(status)`
 - `WTERMSIG(status)` This macro should be employed only if `WIFSIGNALED` returned true.
 - `WCOREDUMP(status)`
 - `WIFSTOPPED(status)`
 - `WSTOPSIG(status)`
 - `WIFCONTINUED(status)`

Multiprocessing – Google Chrome

- Many web browsers ran (in the past) as single process (some still do) → If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multi-process with 3 different types of processes:
 - **Browser** process manages UI, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened → Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in

