


B+ Trees

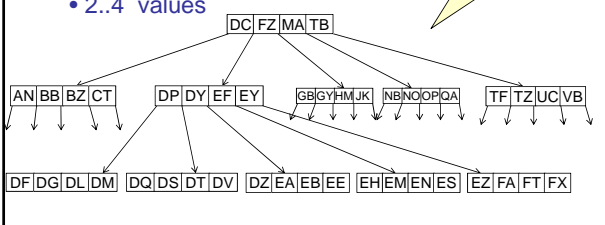


Victoria
UNIVERSITY OF VICTORIA
To Whom It May Concern
CAPITAL CITY UNIVERSITY

B Trees

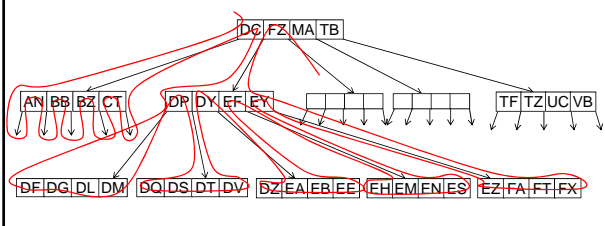
- Example of a B tree with $m = 5$
 - nodes contain
 - 3.5 children
 - 2.4 values

Maximally full depth 3 tree
(Can't fit all the nodes on slide)



Traversing a B Tree

- Listing all the items in order is a bit messy:
 - Have to constantly return to higher nodes.




Sets versus Key-Value pairs

What are the items we are storing?

- Set:
 - B-Tree contains a set of values
 - Only need to store the values you are searching on.
- Map / Dictionary / database table
 - B-Tree contains a set of key-value pairs
 - Search down the tree governed only by the keys
 - Need to store each value with its key
 - ⇒ B-tree nodes can't have as many keys
 - ⇒ lower branching factor
 - ⇒ deeper trees
- If node is of fixed capacity and value is large (eg, a whole record from a database table) then may only fit a very few key-value pairs in a node.

B+ Trees

- The most commonly used variant of B Trees.
- Intended for storing key-value (or key-record) pairs.
- Leaves contain key-value pairs,
- Keys are repeated in the internal nodes.
 - Internal nodes contain keys, but no value
- Leaves are linked to enable traversal



B+ Trees: Leaves

- Each leaf node contains
 - between $\lceil \text{max}_L / 2 \rceil$ and max_L key-value pairs,
 - a link to the next leaf in the tree
- $K_0-V_0, K_1-V_1, K_2-V_2, \dots, K_{\text{leaf.size}-1}-V_{\text{leaf.size}-1}$
- For each key K_i in the leaf :
 - $K_i < K_{i+1}$
 - The value might be either
 - the actual associated value (if it is small)
 - the index of a data block where value can be found (maybe in another file)

B+ Trees: Internal Nodes

Except root may have fewer

- Each internal node contains
 - between $\lfloor \max_N / 2 \rfloor$ and \max_N keys, and
 - up to $\max_N + 1$ child node indexes


$C_0, K_1, C_1, K_2, C_2, \dots, K_{\text{node.size}}, C_{\text{node.size-1}}, C_{\text{node.size}}$

- Branching factor = $\max_N + 1$

- For each key X in the subtree at C_i :


$$K_i \leq X < K_{i+1}$$

- K_i will be the leftmost key in the subtree at C_i ie, the first key in leftmost leaf of C_i



B+ example

- 3-degree, Add in order: M H T S R Q B A F D Z



B+ Tree: Find


To find value associated with a key:

Find(key):

if root is empty return null
else return Find(key, root)

Find(key, node):

if node is a leaf
for i from 0 to node.size-1
if key = node.keys[i] return node.values[i]
return null
if node is an internal node
for i from 1 to node.size
if key < node.keys[i] return Find(key, getNode(node.child[i-1]))
return Find(key, getNode(node.child[node.size]))



Could use binary search


$K_0-V_0, K_1-V_1, K_2-V_2, \dots, K_{\text{leaf.size-1}}-V_{\text{leaf.size-1}}$

B+ Tree Add

- Find leaf node where item belongs
- Insert in leaf .
- If node too full, split, and promote middle key up to parent, middle key also go to the right
- If root split, create new root containing promoted key

Splitting a B+-Tree Leaf


- If a leaf overflows:
 - The left most m keys are left in the node,
 - The right most $m + 1$ keys are moved to a new node,
 - The $(m + 1)$ -st key is propagated to the parent node




- A non leaf node splits as in a common B-tree
- The right sub-tree of each non leaf node contains greater or equal key values

B+-Tree Insertion Example

- Given a B+-tree

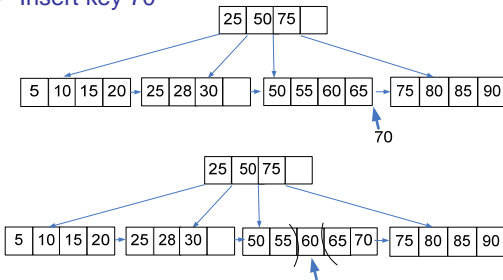


- Insert key 28




B+-Tree Insertion Example (cont.)

- Insert key 70




B+-Tree Insertion Example (cont.)

- The middle key of 60 is placed in the node between 50 and 75




- Insert 95



B+-Tree Insertion Example (cont.)

- Split the leaf and promote the middle key to the parent node

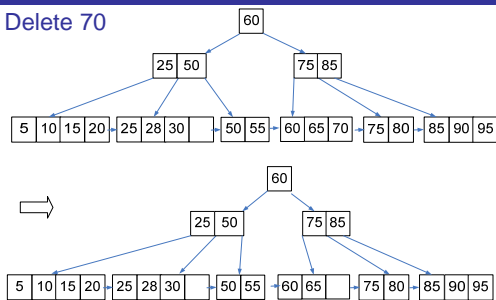


B⁺-Tree Deletion

- When a record is deleted from a B⁺-tree it is always removed from the leaf level
- If the deletion of the key does not cause the leaf underflow
 - Delete the key from the leaf
 - If the key of the deleted record appears in an index node, use the next key to replace it
- If deletion causes the leaf and the corresponding index node underflow
 - Redistribute, if there is a sibling with more than m keys
 - Merge, if there is no sibling with more than m keys
 - Adjust the index node to reflect the change

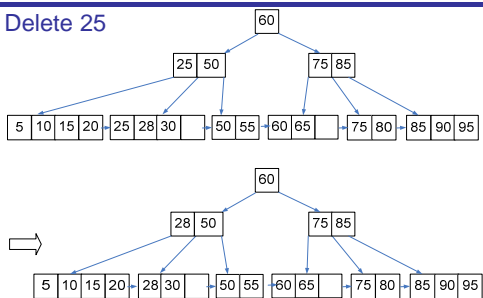
B⁺-Tree Deletion Example

- Delete 70



B⁺-Tree Deletion Example (cont.)

- Delete 25



B+-Tree Deletion Example (cont.)

- Delete 60

