


Indexing Large Data



Introduction

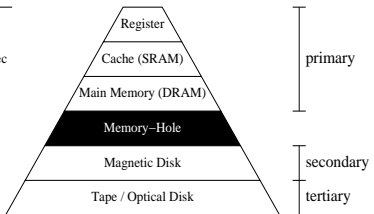
- Files, file structures, DB files, indexes
- B-trees and B+-trees

– Reference Book (in VUW Library):
Fundamentals of Database Systems
 by Elmasri & Navathe (2011),
 (Chapters 16 and 17 only!)

The Memory Hierarchy

- Memory in a computer system is arranged in a hierarchy:


Capacity	Access Time
Byte – KByte	1 nsec
KByte – MByte	2 nsec – 10 nsec
MByte – GByte	50 nsec
GByte	10 msec
GByte – PByte	50 msec – sec



The Memory Hierarchy

- Registers
 - in CPU.
- Memory caches
 - Very fast, small
 - Copy of bits of primary memory
- Primary memory:
 - Fast, large
 - Volatile: lost in power outage
- Secondary Storage (hard disks)
 - Slow, very large
 - Persistent
 - Data cannot be manipulated directly,
 - data must be copied to main memory,
 - modified
 - written back to disk
- Need to know how files are organised on disk

Disk Storage Devices



- Disk controller: interfaces disk drive to computer
- Implements commands to read or write a sector

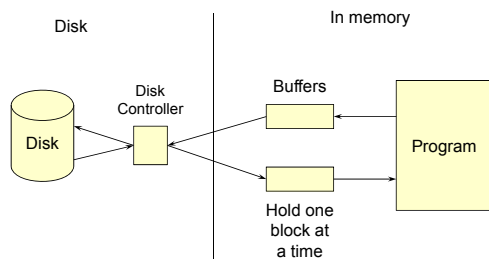
Disk Organisation

Data on a disk is organised into chunks:

- Sectors:
 - physical organisation.
 - hardware disk operations work on sector at a time
 - traditionally: 512 bytes
 - modern disks: 4096 bytes
- Blocks:
 - logical organisation
 - operating system retrieves and writes a block at a time
 - 512 bytes to 8192 bytes,

⇒ For all efficient file operations, minimise *block* accesses


Reading/Writing from files



File Organisation

Files typically much larger than blocks

- A file must be stored in a collection of blocks on disk
 - Must organise the file data into blocks
 - Must record the collection of blocks associated with the file
 - Must enable access to the blocks (sequential or random access)



Files of records

- File may be a sequence of records (especially for DB files)
 - record = logical chunk of information
 - eg a row from a DataBase table
 - an entry in a file system description
 - ...
 - May have several records per block
 - "blocking factor" = # records per block
 - can calculate block number from record number (& vice versa) (if records all the same size)
 - # records = # blocks × bf

Using B+ trees for organising Data




- B+ tree is an efficient index for very large data sets
- The B+ tree must be stored on disk (ie, in a file)
 - \Rightarrow costly actions are accessing data from disk
- Retrieval action in the B+ tree is accessing a node
 - \Rightarrow want to make one node only require one block access
 - \Rightarrow want each of the nodes to be as big as possible \Rightarrow fill the block

B+ tree in a file:

- one node (internal or leaf) per block.
- links to other nodes = index of block in file.
- need some header info.
- need to store keys and values as bytes.

Implementing B+ Tree in a File

- To store a B Tree in a block-structured file
 - Need a block for each node of the tree
 - Can refer to blocks by their index in the file.
 - Need an initial block (first block in file) with meta data:
 - index of the root block
 - number of items
 - information about the item sizes and types?
 - Need a block for each internal node
 - type
 - number of items
 - child key child key... key child
 - Need a block for each leaf node
 - type
 - number of items
 - link to next
 - key-value key-value key-value



index of block containing child node

Cost of B+ tree

- If the block size is 1024 bytes, how big can the nodes be?
- Node requires
 - some header information
 - leaf node or internal node
 - number of items in node,
 - internal node:
 - $m_N \times$ key size
 - $m_N + 1 \times$ pointer size
 - leaf node
 - $m_L \times$ item size
 - pointer to next leaf
- How big is an item?
- How big is a pointer?

Must specify which node, ie which block of the file

Leaf nodes could hold more values than internal nodes!

Cost of B+ tree

- If block has 1024 bytes
- each node has header \Rightarrow 5 bytes
- If key is a string of up to 10 characters \Rightarrow 10 bytes
- if value is a string of up to 20 characters \Rightarrow 20 bytes
- If child pointer is an int \Rightarrow 4 bytes
- Internal node (m_N keys, m_N+1 child pointer)
 - size = 5 + (10 + 4) m_N + 4
 - $\Rightarrow m_N \leq (1024 - 9) / 14 = 72.5 \Rightarrow 72$ keys in internal nodes
- Leaf node (with pointer to next)
 - size = 5 + 4 + (10 + 20) m_L
 - $\Rightarrow m_L \leq (1024 - 9) / 30 = 33.8 \Rightarrow 33$ key-value pairs in leaves

B+ Tree: Find


To find value associated with a key:

Find(key):


```
if root is empty return null
else return Find(key, root)
```

Find(key, node):

```
if node is a leaf
  for i from 0 to node.size-1
    if key = node.keys[i] return node.values[i]
  return null
if node is an internal node
  for i from 1 to node.size
    if key < node.keys[i] return Find(key, getNode(node.child[i-1]))
  return Find(key, getNode(child[node.size]))
```



Could use binary search



B+ Tree Add (1)

Add(key, value):

```
if root is empty
  create new leaf, add key-value,
  root ← leaf
```

else

```
(newKey, rightChild) ← Add(key, value, root)
```

```
if (newKey, rightChild) ≠ null
```

```
  node ← create new internal node
  node.size ← 1
  node.child[0] ← root
  node.keys[1] ← newKey
  node.child[1] ← rightChild
  root ← node
```

If root was full:
returns new key
and new leaf node,

Make a new
root node

B+ Tree Add (2)

Add(key, value, node):

```

if node is a leaf
  if node.size < maxLeafKeys
    insert key and value into leaf in correct place
    return null
  else
    return SplitLeaf(key, value, node)

```

$K_0-V_0, K_1-V_1, K_2-V_2, \dots, K_{size-1}-V_{size-1}$


Returns new key
and new leaf node,

if node is an internal node

```

for i from 1 to node.size
  if key < node.keys[i]
    (k, rc) ← Add(key, value, node.child[i-1])
    if (k, rc) = null return null
    else return dealWithPromote(k, rc, node)
(k, rc) ← Add(key, value, node.child[node.size])
if (k, rc) = null return null
else return dealWithPromote(k, rc, node)

```



Inserts new
key and child
into node,

B+ Tree Add (3)

SplitLeaf(key, value, node):


```

insert key and value into leaf in correct place (spilling over end)
sibling ← create new leaf
mid ← ⌊(node.size+1)/2⌋
move keys and values from mid ... size out of node into sibling.
sibling.next ← node.next node.next ← sibling
return (sibling.keys[0], sibling)

```

Could make the array one
larger than necessary to
give room for this.

$\lfloor (max_L+2)/2 \rfloor$ since size is now max_L+1



B+ Tree Add (4)

DealWithPromote(newKey, rightChild, node):

```

if (newKey, rightChild) = null return null
if newKey > node.keys[node.size]
  insert newKey at node.keys[node.size+1]
  insert rightChild at node.child[node.size+1]

```

```

else for i from 1 to node.size
  if newKey < node.keys[i]
    insert newKey at node.keys[i]
    insert rightChild at node.child[i]

```


if size ≤ maxNodeKeys return null

```

sibling ← create new node
mid ← ⌊size/2⌋ + 1
move node.keys[mid+1... node.size] to sibling.node[1... node.size-mid]
move node.child[mid ... node.size] to sibling.child[0 ... node.size-mid]
promoteKey ← node.keys[mid]
remove node.keys[mid]
return (promoteKey, sibling)

```

Nothing was promoted



No need to promote further
Node is overfull:
Have to split and promote